

Scikit learn

```
In [3]: 1 import pandas as pd
2 import seaborn as sns
3 import matplotlib.pyplot as plt
4 from sklearn.ensemble import RandomForestClassifier
5 from sklearn.svm import SVC
6 from sklearn import svm
7 from sklearn.neural_network import MLPClassifier
8 from sklearn.linear_model import SGDClassifier
9 from sklearn.metrics import confusion_matrix, classification_report
10 from sklearn.preprocessing import StandardScaler, LabelEncoder
11 from sklearn.model_selection import train_test_split
12 %matplotlib inline
```

```
In [69]: 1 df=pd.read_csv(r"C:\Users\user\Downloads\winequality-red.csv")
```

```
In [70]: 1 df.head(10)
```

Out[70]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4
5	7.4	0.66	0.00	1.8	0.075	13.0	40.0	0.9978	3.51	0.56	9.4
6	7.9	0.60	0.06	1.6	0.069	15.0	59.0	0.9964	3.30	0.46	9.4
7	7.3	0.65	0.00	1.2	0.065	15.0	21.0	0.9946	3.39	0.47	10.0
8	7.8	0.58	0.02	2.0	0.073	9.0	18.0	0.9968	3.36	0.57	9.5
9	7.5	0.50	0.36	6.1	0.071	17.0	102.0	0.9978	3.35	0.80	10.5

In [71]:

```
1 df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide    1599 non-null   float64
6   total sulfur dioxide   1599 non-null   float64
7   density               1599 non-null   float64
8   pH                    1599 non-null   float64
9   sulphates              1599 non-null   float64
10  alcohol                1599 non-null   float64
11  quality                1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

In [72]:

```
1 df.isnull().sum()
```

```
Out[72]: fixed acidity          0
volatile acidity          0
citric acid               0
residual sugar            0
chlorides                 0
free sulfur dioxide       0
total sulfur dioxide      0
density                   0
pH                        0
sulphates                 0
alcohol                   0
quality                   0
dtype: int64
```

In [73]:

```
1 df.head()
```

Out[73]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4

```
In [86]: 1 #preprocessing data
2 bins=(2,6.5,8)
3 group_names=["good","bad"]
4 df["quality"]=pd.cut(df["quality"],bins=bins,labels=group_names)
5 df["quality"].unique()
```

```
Out[86]: [NaN, 'good']
Categories (2, object): ['good' < 'bad']
```

```
In [87]: 1 label_quality=LabelEncoder()
```

```
In [88]: 1 df["quality"]=label_quality.fit_transform(df["quality"])
```

```
In [89]: 1 df.head(10)
```

```
Out[89]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4
5	7.4	0.66	0.00	1.8	0.075	13.0	40.0	0.9978	3.51	0.56	9.4
6	7.9	0.60	0.06	1.6	0.069	15.0	59.0	0.9964	3.30	0.46	9.4
7	7.3	0.65	0.00	1.2	0.065	15.0	21.0	0.9946	3.39	0.47	10.0
8	7.8	0.58	0.02	2.0	0.073	9.0	18.0	0.9968	3.36	0.57	9.5
9	7.5	0.50	0.36	6.1	0.071	17.0	102.0	0.9978	3.35	0.80	10.5


```
In [90]: 1 df["quality"].value_counts()
```

```
Out[90]: 0    855
1    744
Name: quality, dtype: int64
```

```
In [91]: 1 sns.countplot(df["quality"])
```

C:\Users\user\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(

```
Out[91]: <AxesSubplot:xlabel='quality', ylabel='count'>
```



```
In [97]: 1 #seperating dataset as response variable and feature variable  
2 x=df.drop("quality",axis=1)  
3 y=df["quality"]
```

```
In [102]: 1 #train and testing the splitting of data  
2 x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.2,random
```

```
In [103]: 1 #applying standard scaler to get optimized results  
2 sc=StandardScaler()  
3 x_train=sc.fit_transform(x_train)  
4 x_test=sc.transform(x_test)
```

```
In [106]: 1 x_train[:10]
```

```
Out[106]: array([[ 0.21833164,  0.88971201,  0.19209222,  0.30972563, -0.04964208,
  0.69100692,  1.04293362,  1.84669643,  1.09349989,  0.45822284,
  1.12317723],
 [-1.29016623, -1.78878251,  0.65275338, -0.80507963, -0.45521361,
  2.38847304,  3.59387025, -3.00449133, -0.40043872, -0.40119696,
  1.40827174],
 [ 1.49475291, -0.78434707,  1.01104539, -0.52637831,  0.59927236,
 -0.95796016, -0.99174203,  0.76865471, -0.07566946,  0.51551749,
 -0.58738978],
 [ 0.27635078,  0.86181102, -0.06383064, -0.66572897, -0.00908493,
  0.01202048, -0.71842739,  0.08948842,  0.05423824, -1.08873281,
 -0.96751578],
 [ 0.04427419,  2.81487994, -0.62686095,  2.39998549, -0.31326357,
 -0.47296984,  0.2229897 ,  1.1998714 ,  0.37900751, -0.9741435 ,
 -0.49235828],
 [-0.07176411, -0.78434707,  1.11341454, -0.17800167,  0.21397941,
  3.01896045,  2.62208486,  0.60694845,  0.44396136,  1.89058918,
 -0.58738978],
 [-1.17412793,  0.10848444, -0.62686095, -0.52637831, -0.23214927,
  0.98200112, -0.35400787, -1.95879086,  0.05423824,  0.91658007,
  1.12317723],
 [-0.1878024 , -0.17052541,  0.60156881,  0.03102432, -0.13075639,
 -0.37597178, -0.01995665,  0.93036097,  0.76873063, -0.229313 ,
  0.26789373],
 [-0.07176411,  0.61070216, -0.01264607, -0.38702766,  0.13286511,
 -1.05495822,  0.92146044,  0.37516948, -1.17988496, -0.229313 ,
 -1.25261029],
 [ 1.8428678 , -1.95618842,  1.21578369,  1.00647892,  0.31537229,
 -1.15195628, -0.71842739,  1.52328391, -0.20557717,  1.77599987,
 -0.30229528]])
```

Random forest classifier

```
In [109]: 1 rfc=RandomForestClassifier(n_estimators=200)
          2 rfc.fit(x_train,y_train)
          3 pred_rfc = rfc.predict(x_test)
```

```
In [112]: 1 pred_rfc[:20]
```

```
Out[112]: array([1, 1, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1])
```

```
In [116]: 1 #model performing
          2 print(classification_report(y_test,pred_rfc))
```

	precision	recall	f1-score	support
0	0.81	0.81	0.81	179
1	0.76	0.76	0.76	141
accuracy			0.79	320
macro avg	0.78	0.78	0.78	320
weighted avg	0.79	0.79	0.79	320

Svm classifier

```
In [117]: 1 clf=svm.SVC()
          2 clf.fit(x_train,y_train)
          3 pred_clf=clf.predict(x_test)
```

```
In [118]: 1 print(classification_report(y_test,pred_clf))
          2 print(confusion_matrix(y_test,pred_clf))
```

	precision	recall	f1-score	support
0	0.81	0.77	0.79	179
1	0.73	0.77	0.75	141
accuracy			0.77	320
macro avg	0.77	0.77	0.77	320
weighted avg	0.77	0.77	0.77	320

```
[[138 41]
 [ 32 109]]
```

Neural network

```
In [121]: 1 mlpc=MLPClassifier(hidden_layer_sizes=(11,11,11),max_iter=500)
          2 mlpc.fit(x_train,y_train)
          3 pred_mlpc=mlpc.predict(x_test)
```

C:\Users\user\anaconda3\lib\site-packages\sklearn\neural_network_multilayer_perceptron.py:692: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (500) reached and the optimization hasn't converged yet.

```
warnings.warn(
```

```
In [122]: 1 print(classification_report(y_test,pred_mlpc))
          2 print(confusion_matrix(y_test,pred_mlpc))
```

```

              precision    recall  f1-score   support

     0       0.78       0.78       0.78        179
     1       0.72       0.72       0.72        141

 accuracy          0.75          0.75          0.75          320
 macro avg          0.75          0.75          0.75          320
 weighted avg          0.75          0.75          0.75          320

[[139  40]
 [ 40 101]]
```

```
In [123]: 1 from sklearn.metrics import accuracy_score
          2 cm=accuracy_score(y_test,pred_rfc)
          3 cm
```

Out[123]: 0.7875

```
In [124]: 1 df.head()
```

Out[124]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4